

## Step 3: data Cleaning

### 1: Drop all punctuation and transform all words into lower case.

```
spooky_wrd<-unnest_tokens(spooky,word,text)
head(spooky_wrd)
```

```
##           id author      word
## 1   id26305    EAP      this
## 1.1 id26305    EAP    process
## 1.2 id26305    EAP   however
## 1.3 id26305    EAP  afforded
## 1.4 id26305    EAP        me
## 1.5 id26305    EAP        no
```

### 2: Bi-grams, n-grams

If we wanna get relationships between words, we use n-grams. So far we've considered words as individual units, and considered their relationships to sentiments or to documents. However, many interesting text analyses are based on the relationships between words, whether examining which words tend to follow others immediately. we'll explore some of the methods tidytext offers for calculating and visualizing relationships between words in your text dataset. This includes the token = "ngrams" argument, which tokenizes by pairs of adjacent words rather than by individual ones. We'll also introduce two new packages: ggraph, which extends ggplot2 to construct network plots, and widyr, which calculates pairwise correlations and distances within a tidy data frame. Together these expand our toolbox for exploring text within the tidy data framework.

#### (1): Tokenizing by n-gram

We've been using the unnest\_tokens function to tokenize by word, or sometimes by sentence, which is useful for the kinds of sentiment and frequency analyses we've been doing so far. But we can also use the function to tokenize into consecutive sequences of words, called n-grams. By seeing how often word X is followed by word Y, we can then build a model of the relationships between them. We do this by adding the token = "ngrams" option to unnest\_tokens(), and setting n to the number of words we wish to capture in each n-gram. When we set n to 2, we are examining pairs of two consecutive words, often called "bigrams"

```
# Make a table with one word per row and remove `stop words` (i.e. the common words).
bigrams<-unnest_tokens(spooky,bigram, text, token = "ngrams", n = 2)
head(bigrams)
```

```
##           id author      bigram
## 1 id00001    MWS   idris was
## 2 id00001    MWS    was well
## 3 id00001    MWS well content
## 4 id00001    MWS content with
## 5 id00001    MWS   with this
## 6 id00001    MWS this resolve
```

```
bigrams_HPL<-unnest_tokens(spooky[spooky$author=='HPL'],bigram, text, token = "ngrams", n = 2)
head(bigrams_HPL)
```

```
##           id author      bigram
## 1 id00002    HPL      i was
## 2 id00002    HPL    was faint
```

```
## 3 id00002    HPL    faint even
## 4 id00002    HPL even fainter
## 5 id00002    HPL fainter than
## 6 id00002    HPL      than the
```

```
bigrams_MWS<-unnest_tokens(spooky[spooky$author=='MWS'],bigram, text, token = "ngrams", n = 2)
head(bigrams_MWS)
```

```
##          id author      bigram
## 1 id00001    MWS    idris was
## 2 id00001    MWS      was well
## 3 id00001    MWS well content
## 4 id00001    MWS content with
## 5 id00001    MWS   with this
## 6 id00001    MWS this resolve
```

```
bigrams_EAP<-unnest_tokens(spooky[spooky$author=='EAP'],bigram, text, token = "ngrams", n = 2)
head(bigrams_EAP)
```

```
##          id author      bigram
## 1 id00003    EAP above all
## 2 id00003    EAP    all i
## 3 id00003    EAP    i burn
## 4 id00003    EAP  burn to
## 5 id00003    EAP   to know
## 6 id00003    EAP know the
```

This data structure is still a variation of the tidy text format. It is structured as one-token-per-row (with extra metadata, such as author, still preserved), but each token now represents a bigram.